

```

import pandas as pd
import numpy as np
import pickle
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import sklearn

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import GradientBoostingClassifier, RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import RandomizedSearchCV
import imblearn

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, f1_score

```

```

dataset = pd.read_csv("flightdata.csv")

dataset.head()

```

YEAR	QUARTER	MONTH	DAY_OF_MONTH	DAY_OF_WEEK	UNIQUE_CARRIER	TAIL_NUM	FL_NUM	ORIGIN_AIRPORT_ID	ORIGIN	CRS_ARR_TIME	ARR_TIME	ARR_DELAY	ARR_DEL15	CANCELLED	DIVERTED	CRS_ELAPSED_TIME	ACTUAL_ELAPSED_TIME	DISTANCE	Unnamed: 25
2016	1	1	1	5	DL	N333DN	1399	10397	ATL	2143	2102.0	-41.0	0.0	0.0	0.0	338.0	295.0	2182.0	NaN
2016	1	1	1	5	DL	N364DN	1476	11433	DTW	1435	1439.0	4.0	0.0	0.0	0.0	110.0	115.0	528.0	NaN
2016	1	1	1	5	DL	N313DN	1987	10397	ATL	1215	1142.0	-33.0	0.0	0.0	0.0	331.0	300.0	2182.0	NaN
2016	1	1	1	5	DL	N557NW	1768	14747	SEA	1335	1345.0	10.0	0.0	0.0	0.0	196.0	205.0	1399.0	NaN
2016	1	1	1	5	DL	N334DN	1023	14747	SEA	607	615.0	8.0	0.0	0.0	0.0	247.0	259.0	1927.0	NaN

rows x 26 columns

```

dataset.info()

```

```

1721
...
Output exceeds the size limit. Open the full output data in a text editor
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 11231 entries, 0 to 11230
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  ---                ---
0   YEAR                  11231 non-null  int64
1   QUARTER               11231 non-null  int64
2   MONTH                11231 non-null  int64
3   DAY_OF_MONTH          11231 non-null  int64
4   DAY_OF_WEEK           11231 non-null  int64
5   UNIQUE_CARRIER       11231 non-null  object
6   TAIL_NUM              11231 non-null  object
7   FL_NUM               11231 non-null  int64
8   ORIGIN_AIRPORT_ID     11231 non-null  int64
9   ORIGIN                11231 non-null  object
10  DEST_AIRPORT_ID       11231 non-null  int64
11  DEST                  11231 non-null  object
12  CRS_DEP_TIME          11231 non-null  int64
13  DEP_TIME              11124 non-null  float64
14  DEP_DELAY             11124 non-null  float64
15  DEP_DEL15             11124 non-null  float64
16  CRS_ARR_TIME          11231 non-null  int64
17  ARR_TIME              11116 non-null  float64
18  ARR_DELAY             11043 non-null  float64
19  ARR_DEL15             11043 non-null  float64

```

```
> dataset = dataset.drop('Unnamed: 25', axis=1)
dataset.isnull().sum()

[10]

...
YEAR          0
QUARTER        0
MONTH          0
DAY_OF_MONTH   0
DAY_OF_WEEK    0
UNIQUE_CARRIER 0
TAIL_NUM       0
FL_NUM         0
ORIGIN_AIRPORT_ID 0
ORIGIN         0
DEST_AIRPORT_ID 0
DEST           0
CRS_DEP_TIME    0
DEP_TIME       107
DEP_DELAY       107
DEP_DEL15       107
CRS_ARR_TIME    0
ARR_TIME       115
ARR_DELAY       188
ARR_DEL15       188
CANCELLED       0
DIVERTED        0
CRS_ELAPSED_TIME 0
ACTUAL_ELAPSED_TIME 188
DISTANCE        0
dtype: int64
```

```
#filter the dataset to eliminate columns that aren't relevant to a predictive model.
dataset = dataset[['FL_NUM', 'MONTH', 'DAY_OF_MONTH', 'DAY_OF_WEEK', 'ORIGIN', 'DEST', 'CRS_ARR_TIME', 'DEP_DEL15', 'ARR_DEL15']]
dataset.isnull().sum()

FL_NUM      0
MONTH        0
DAY_OF_MONTH 0
DAY_OF_WEEK  0
ORIGIN       0
DEST         0
CRS_ARR_TIME 0
DEP_DEL15    107
ARR_DEL15    188
dtype: int64

dataset[dataset.isnull().any(axis=1)].head(10)
```

	FL_NUM	MONTH	DAY_OF_MONTH	DAY_OF_WEEK	ORIGIN	DEST	CRS_ARR_TIME	DEP_DEL15	ARR_DEL15
177	2834	1	9	6	MSP	SEA	852	0.0	NaN
179	86	1	10	7	MSP	DTW	1632	NaN	NaN
184	557	1	10	7	MSP	DTW	912	0.0	NaN
210	1096	1	10	7	DTW	MSP	1303	NaN	NaN
478	1542	1	22	5	SEA	JFK	723	NaN	NaN
481	1795	1	22	5	ATL	JFK	2014	NaN	NaN
491	2312	1	22	5	MSP	JFK	2149	NaN	NaN
499	423	1	23	6	JFK	ATL	1600	NaN	NaN
500	425	1	23	6	JFK	ATL	1827	NaN	NaN
501	427	1	23	6	JFK	SEA	1053	NaN	NaN

```
dataset['DEP_DEL15'].mode()

0    0.0
dtype: float64

#replace the missing values with 1s.
dataset = dataset.fillna({'ARR_DEL15': 1})
dataset = dataset.fillna({'DEP_DEL15': 0})
dataset.iloc[177:185]
```

	FL_NUM	MONTH	DAY_OF_MONTH	DAY_OF_WEEK	ORIGIN	DEST	CRS_ARR_TIME	DEP_DEL15	ARR_DEL15
177	2834	1	9	6	MSP	SEA	852	0.0	1.0
178	2839	1	9	6	DTW	JFK	1724	0.0	0.0
179	86	1	10	7	MSP	DTW	1632	0.0	1.0
180	87	1	10	7	DTW	MSP	1649	1.0	0.0
181	423	1	10	7	JFK	ATL	1600	0.0	0.0
182	440	1	10	7	JFK	ATL	849	0.0	0.0
183	485	1	10	7	JFK	SEA	1945	1.0	0.0
184	557	1	10	7	MSP	DTW	912	0.0	1.0

```
import math

for index, row in dataset.iterrows():
    dataset.loc[index, 'CRS_ARR_TIME'] = math.floor(row['CRS_ARR_TIME'] / 100)
dataset.head()
```

FL_NUM	MONTH	DAY_OF_MONTH	DAY_OF_WEEK	ORIGIN	DEST	CRS_ARR_TIME	DEP_DEL15	ARR_DEL15
1399	1	1	5	ATL	SEA	21	0.0	0.0
1476	1	1	5	DTW	MSP	14	0.0	0.0
1597	1	1	5	ATL	SEA	12	0.0	0.0
1768	1	1	5	SEA	MSP	13	0.0	0.0
1823	1	1	5	SEA	DTW	6	0.0	0.0

```
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
dataset['DEST'] = le.fit_transform(dataset['DEST'])
dataset['ORIGIN'] = le.fit_transform(dataset['ORIGIN'])
```

```
dataset.head(5)
```

FL_NUM	MONTH	DAY_OF_MONTH	DAY_OF_WEEK	ORIGIN	DEST	CRS_ARR_TIME	DEP_DEL15	ARR_DEL15
1399	1	1	5	0	4	21	0.0	0.0
1476	1	1	5	1	3	14	0.0	0.0
1597	1	1	5	0	4	12	0.0	0.0
1768	1	1	5	4	3	13	0.0	0.0
1823	1	1	5	4	1	6	0.0	0.0

```
dataset['ORIGIN'].unique()
```

```
array([0, 1, 4, 3, 2])
```

```
dataset = pd.get_dummies(dataset, columns=['ORIGIN', 'DEST'])
dataset.head()
```

```
x = dataset.iloc[:, 0:8].values
y = dataset.iloc[:, 8:9].values
```

```
x
```

```
array([[1.399e+03, 1.000e+00, 1.000e+00, ..., 4.000e+00, 2.100e+01,
        0.000e+00],
       [1.476e+03, 1.000e+00, 1.000e+00, ..., 3.000e+00, 1.400e+01,
        0.000e+00],
       [1.597e+03, 1.000e+00, 1.000e+00, ..., 4.000e+00, 1.200e+01,
        0.000e+00],
       ...,
       [1.823e+03, 1.200e+01, 3.000e+01, ..., 4.000e+00, 2.200e+01,
        0.000e+00],
       [1.901e+03, 1.200e+01, 3.000e+01, ..., 4.000e+00, 1.800e+01,
        0.000e+00],
       [2.005e+03, 1.200e+01, 3.000e+01, ..., 1.000e+00, 9.000e+00,
        0.000e+00]])
```

```

from sklearn.preprocessing import OneHotEncoder
oh = OneHotEncoder()
z=oh.fit_transform(x[:,4:5]).toarray()
t=oh.fit_transform(x[:,5:6]).toarray()
#x=np.delete(x,[4,7],axis=1)

```

z

```

array([[1., 0., 0., 0., 0.],
       [0., 1., 0., 0., 0.],
       [1., 0., 0., 0., 0.],
       ...,
       [0., 1., 0., 0., 0.],
       [1., 0., 0., 0., 0.],
       [1., 0., 0., 0., 0.]])

```

t

```

array([[0., 0., 0., 0., 1.],
       [0., 0., 0., 1., 0.],
       [0., 0., 0., 0., 1.],
       ...,
       [0., 0., 0., 0., 1.],
       [0., 0., 0., 0., 1.],
       [0., 1., 0., 0., 0.]])

```

```

x=np.delete(x,[4,5],axis=1)

```

flight\_data.describe()

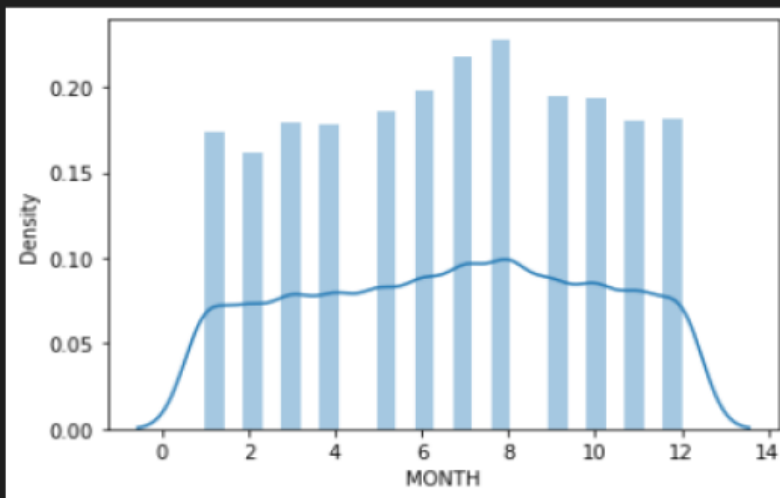
	YEAR	QUARTER	MONTH	DAY_OF_MONTH	DAY_OF_WEEK	FL_NUM	ORIGIN_AIRPORT_ID	DEST_AIRPORT_ID	CRS_DEP_TIME	DEP_TIME	CRS_ARR_TIME	ARR_TIME	ARR_DELAY	ARR_DEL15	CANCELLED	DIVERTED	CRS_ELAPSED_TIME	ACTUAL_ELAPSED_TIME	DISTANCE	Unnamed: 25
count	11231.0	11231.000000	11231.000000	11231.000000	11231.000000	11231.000000	11231.000000	11231.000000	11231.000000	11124.000000	11231.000000	11116.000000	11043.000000	11043.000000	11231.000000	11231.000000	11231.000000	11043.000000	11231.000000	0.0
mean	2016.0	2.544475	6.632073	15.700758	3.062109	1334.325917	12334.516595	12302.274538	1320.708236	1327.180410	1537.312795	1533.378400	-2.579123	0.124513	0.010150	0.006580	100.652124	179.661233	1161.031965	NaN
std	0.0	1.090701	3.354678	8.782056	1.995257	811.875227	1595.029510	1601.508550	480.737845	500.306462	502.512484	512.536041	39.232521	0.330181	0.100241	0.008908	78.318617	77.940399	643.683379	NaN
min	2016.0	1.000000	1.000000	1.000000	1.000000	7.000000	10397.000000	10397.000000	10.000000	1.000000	2.000000	1.000000	-57.000000	0.000000	0.000000	0.000000	93.000000	75.000000	508.000000	NaN
25%	2016.0	2.000000	4.000000	8.000000	2.000000	624.000000	10397.000000	10397.000000	905.000000	905.000000	1130.000000	1135.000000	-19.000000	0.000000	0.000000	0.000000	127.000000	117.000000	594.000000	NaN
50%	2016.0	3.000000	7.000000	16.000000	4.000000	1267.000000	12478.000000	12478.000000	1320.000000	1324.000000	1559.000000	1547.000000	-10.000000	0.000000	0.000000	0.000000	159.000000	148.000000	907.000000	NaN
75%	2016.0	3.000000	9.000000	23.000000	6.000000	2052.000000	13487.000000	13487.000000	1795.000000	1739.000000	1952.000000	1945.000000	1.000000	0.000000	0.000000	0.000000	255.000000	236.000000	1927.000000	NaN
max	2016.0	4.000000	12.000000	31.000000	7.000000	2853.000000	14747.000000	14747.000000	2359.000000	2400.000000	2359.000000	2400.000000	615.000000	1.000000	1.000000	1.000000	397.000000	438.000000	3422.000000	NaN

rows x 22 columns

```
sns.distplot(flight_data.MONTH)
```


C:\Users\Saumya\Anaconda3\lib\site-packages\seaborn\distributions.py:2557:  
figure-level function with similar flexibility) or `histplot` (an axes-level  
warnings.warn(msg, FutureWarning)

<AxesSubplot:xlabel='MONTH', ylabel='Density'>




```
sns.scatterplot(x='ARR_DELAY',y='ARR_DEL15',data=flight_data)
```

<AxesSubplot:xlabel='ARR\_DELAY', ylabel='ARR\_DEL15'>



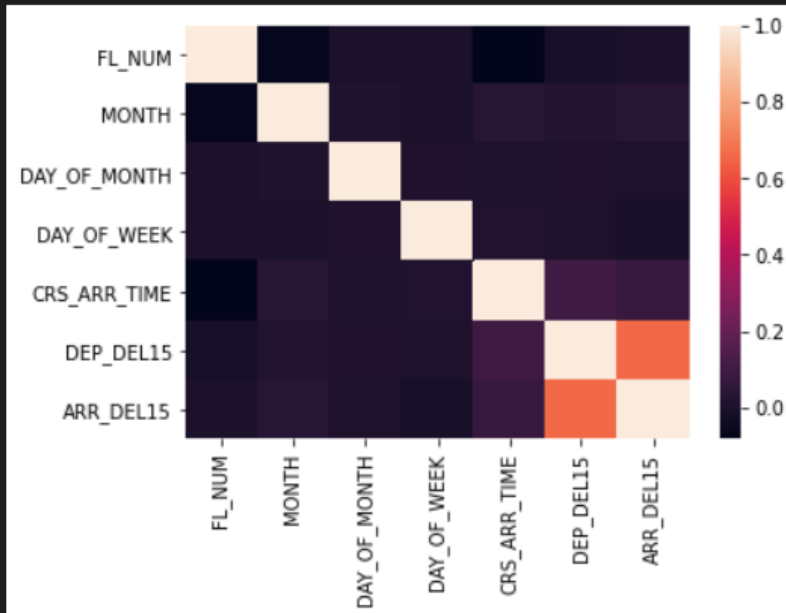
```
sns.catplot(x="ARR_DEL15",y="ARR_DELAY",kind='bar',data=flight_data)
```

<seaborn.axisgrid.FacetGrid at 0x22716099eb0>



```
sns.heatmap(dataset.corr())
```

<AxesSubplot:>



```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_state=0)
```

```
from sklearn.model_selection import train_test_split
train_x, test_x, train_y, test_y = train_test_split(dataset.drop('ARR_DEL15', axis=1), df['ARR_DEL15'], test_size=0.2, random_state=0)
```

```
x_test.shape
```

```
(2247, 16)
```

```
x_train.shape
```

```
(8984, 16)
```

```
y_test.shape
```

```
(2247, 1)
```

[+ Code](#)[+ Markdown](#)

```
y_train.shape
```

```
(8984, 1)
```

▼

```
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
x_train = sc.fit_transform(x_train)
x_test = sc.transform(x_test)
```

```
8]
```



```
3] from sklearn.tree import DecisionTreeClassifier
   classifier = DecisionTreeClassifier(random_state = 0)
   classifier.fit(x_train,y_train)
```

```
DecisionTreeClassifier(random_state=0)
```

```
4] decisiontree = classifier.predict(x_test)
```

```
5] decisiontree
```

```
array([1., 0., 0., ..., 0., 0., 1.])
```

```
6] from sklearn.metrics import accuracy_score
   desacc = accuracy_score(y_test,decisiontree)
```

```
7] from sklearn.ensemble import RandomForestClassifier
   rfc = RandomForestClassifier(n_estimators=10,criterion='entropy')
```

```
8] rfc.fit(x_train,y_train)
```

```
<ipython-input-125-b87bb2ba9825>:1: DataConversionWarning: A column-vector y was passed; you can fix the warning flag with ravel().
```

```
   rfc.fit(x_train,y_train)
```

```
RandomForestClassifier(criterion='entropy', n_estimators=10)
```

```
9] y_predict = rfc.predict(x_test)
```

```
# Importing the Keras libraries and packages
import tensorflow
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
```

```
# Creating ANN skleton view
```

```
classification = Sequential()
classification.add(Dense(30,activation='relu'))
classification.add(Dense(128,activation='relu'))
classification.add(Dense(64,activation='relu'))
classification.add(Dense(32,activation='relu'))
classification.add(Dense(1,activation='sigmoid'))
```

```
# Compiling the ANN model
```

```
classification.compile(optimizer='adam',loss='binary_crossentropy',metrics=['accuracy'])
```

```
# Training the model
```

```
classification.fit(x_train,y_train,batch_size=4,validation_split=0.2,epochs=100)
```

Output exceeds the [size limit](#). Open the full output data [in a text editor](#)

Epoch 1/100

1797/1797 [=====] - 6s 2ms/step - loss: 0.2873 - accuracy: 0.8988 - val\_loss: 0.2722 - val\_accuracy: 0.9071

'''

Epoch 99/100 1797/1797 [=====] - 4s 2ms/step - loss: 0.0586 - accuracy: 0.9789 - val\_loss: 1.1199 - val\_accuracy: 0.8676 Epoch 100/100 1797/1797  
[=====] - 5s 3ms/step - loss: 0.0517 - accuracy: 0.9811 - val\_loss: 1.1271 - val\_accuracy: 0.8648

<tensorflow.python.keras.callbacks.History at 0x22721bdb7c0>

```
## Decision tree
```

```
y_pred = classifier.predict([[129,99,1,0,0,1,0,1,1,1,0,1,1,1,1]])
```

```
print(y_pred)  
(y_pred)
```

```
[0.]
```

```
array([0.])
```

```
## RandomForest
```

```
y_pred = rfc.predict([[129,99,1,0,0,1,0,1,1,1,0,1,1,1,1]])
```

```
print(y_pred)  
(y_pred)
```

```
[0.]
```

```
array([0.])
```

```
classification.save('flight.h5')
```

```
# Testing the model
```

```
y_pred = classification.predict(x_test)
```

```
y_pred
```

```
array([[3.1306639e-01],  
       [4.3961532e-19],  
       [8.1048012e-03],  
       ...,  
       [1.5726548e-10],  
       [3.8635731e-04],  
       [9.994898e-01]], dtype=float32)
```

```
y_pred = (y_pred > 0.5)
y_pred
```

66]

```
.. array([[False],
          [False],
          [False],
          ...,
          [False],
          [False],
          [ True]])
```

```
def predict_exit(sample_value):
    # Convert list to numpy array
    sample_value = np.array(sample_value)

    # Reshape because sample_value contains only 1 record
    sample_value = sample_value.reshape(1, -1)

    # Feature Scaling
    sample_value = sc.transform(sample_value)

    return classifier.predict(sample_value)
```

8]

```
test=classification.predict([[1,1,121.000000,36.0,0,0,1,0,1,1,1,1,1,1,1]])
if test==1:
    print('Prediction: Chance of delay')
else:
    print('Prediction: No chance of delay.')
```

0]

Prediction: No chance of delay.

```
from sklearn import model_selection
from sklearn.neural_network import MLPClassifier
```

```

dfs = []
models = [
    ('RF', RandomForestClassifier()),
    ('DecisionTree', DecisionTreeClassifier()),
    ('ANN', MLPClassifier())
]
results = []
names = []
scoring = ['accuracy', 'precision_weighted', 'recall_weighted', 'f1_weighted', 'roc_auc']
target_names = ['no delay', 'delay']
for name, model in models:
    kfold = model_selection.KFold(n_splits=5, shuffle=True, random_state=90210)
    cv_results = model_selection.cross_validate(model, x_train, y_train, cv=kfold, scoring=scoring)
    clf = model.fit(x_train, y_train)
    y_pred = clf.predict(x_test)
    print(name)
    print(classification_report(y_test, y_pred, target_names=target_names))
    results.append(cv_results)
    names.append(name)
    this_df = pd.DataFrame(cv_results)
    this_df['model'] = name
    dfs.append(this_df)
final = pd.concat(dfs, ignore_index=True)
return final

```

#### RF

	precision	recall	f1-score	support
no delay	0.93	0.96	0.95	1936
delay	0.72	0.58	0.64	311
accuracy			0.91	2247
macro avg	0.82	0.77	0.79	2247
weighted avg	0.90	0.91	0.91	2247

#### DecisionTree

	precision	recall	f1-score	support
no delay	0.93	0.93	0.93	1936
delay	0.56	0.55	0.55	311
accuracy			0.88	2247
macro avg	0.74	0.74	0.74	2247
weighted avg	0.88	0.88	0.88	2247

ANN

	precision	recall	f1-score	support
no delay	0.93	0.96	0.95	1936
delay	0.70	0.58	0.63	311
accuracy			0.91	2247
macro avg	0.82	0.77	0.79	2247
weighted avg	0.90	0.91	0.90	2247

```
# RandomForest Accuracy
print('Training accuracy: ',accuracy_score(y_train,y_predict_train))
print('Testing accuracy: ',accuracy_score(y_test,y_predict))
```

```
Training accuracy:  0.9892030276046304
Testing accuracy:  0.89942145082332
```

```
# Making the Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_predict)
cm
```

```
array([[1874,  62],
       [ 161, 150]], dtype=int64)
```

```
# Accuracy score of desicionTree
```

```
from sklearn.metrics import accuracy_score  
desacc = accuracy_score(y_test,decisiontree)
```

```
desacc
```

```
0.8673787271918113
```

```
from sklearn.metrics import confusion_matrix  
cm = confusion_matrix(y_test,decisiontree)
```

```
cm
```

```
array([[1777, 159],  
       [ 139, 172]], dtype=int64)
```

```
# Calculate the Accuracy of ANN
```

```
from sklearn.metrics import accuracy_score,classification_report  
score = accuracy_score(y_pred,y_test)  
print('The accuracy for ANN model is: {}'.format(score*100))
```

```
The accuracy for ANN model is: 87.2719181130396%
```

```
# Making the Confusion Matrix
```

```
from sklearn.metrics import confusion_matrix  
cm = confusion_matrix(y_test, y_pred)  
cm
```

```
array([[1812, 124],  
       [ 162, 149]], dtype=int64)
```

```
# giving some parameters that can be used in randomized search cv
parameters = {
    'n_estimators' : [1,20,30,55,68,74,90,120,115],
    'criterion':['gini','entropy'],
    'max_features' : ["auto", "sqrt", "log2"],
    'max_depth' : [2,5,8,10], 'verbose' : [1,2,3,4,6,8,9,10]
}
```

```
#performing the randomized cv
RCV = RandomizedSearchCV(estimator=rf,param_distributions=parameters,cv=10,n_iter=4)
```

```
RCV.fit(x_train,y_train)
```

```
building tree 89 of 90
```

```
building tree 90 of 90
```

```
[Parallel(n_jobs=1)]: Done 90 out of 90 | elapsed: 1.1s finished
```

```
RandomizedSearchCV(cv=10, estimator=RandomForestClassifier(), n_iter=4,
    param_distributions={'criterion': ['gini', 'entropy'],
        'max_depth': [2, 5, 8, 10],
        'max_features': ['auto', 'sqrt',
            'log2'],
        'n_estimators': [1, 20, 30, 55, 68, 74,
            90, 120, 115],
        'verbose': [1, 2, 3, 4, 6, 8, 9, 10]})
```

```
#getting the best paarmets from the giving list and best score from them
bt_params = RCV.best_params_
bt_score = RCV.best_score_
```



```
bt_params
```

```
{'verbose': 10,  
 'n_estimators': 90,  
 'max_features': 'log2',  
 'max_depth': 10,  
 'criterion': 'entropy'}
```

```
bt_score
```

```
0.905498809615237
```

```
model = RandomForestClassifier(verbose=10, n_estimators=120, max_features='log2', max_depth=10, criterion='entropy')  
RCV.fit(x_train, y_train)
```

```
RandomizedSearchCV(cv=10, estimator=RandomForestClassifier(), n_iter=4,  
                  param_distributions={'criterion': ['gini', 'entropy'],  
                                       'max_depth': [2, 5, 8, 10],  
                                       'max_features': ['auto', 'sqrt',  
                                                       'log2'],  
                                       'n_estimators': [1, 20, 30, 55, 68, 74,  
                                                       90, 120, 115],  
                                       'verbose': [1, 2, 3, 4, 6, 8, 9, 10]})
```

```
y_predict_rf = RCV.predict(x_test)
```

```
[Parallel(n_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.  
[Parallel(n_jobs=1)]: Done 115 out of 115 | elapsed: 0.0s finished
```

```
RFC=accuracy_score(y_test, y_predict_rf)  
RFC
```

```
0.9096573208722741
```

```
import pickle  
pickle.dump(RCV,open('flight.pkl','wb'))
```